

# How to handle large model output

ACCESS Training Day 2024


What's so hard about processing large datasets?



# What's so hard about processing large datasets?

Common problems

Data is bigger than  
available memory



Takes a really long  
time to process

# What's so hard about processing large datasets?

Common problems

Data is bigger than  
available memory



Takes a really long  
time to process



Helpful strategies

Parallel & distributed  
computing

Read and process data in  
chunks



Lazy loading/evaluation

Out-of-core computation

# What's so hard about processing large datasets?

Common problems

Data is bigger than  
available memory



Takes a really long  
time to process



Helpful strategies

Parallel & distributed  
computing

Read and process data in  
chunks



Lazy loading/evaluation

Out-of-core computation

Useful software




# What is Xarray and what is Dask?

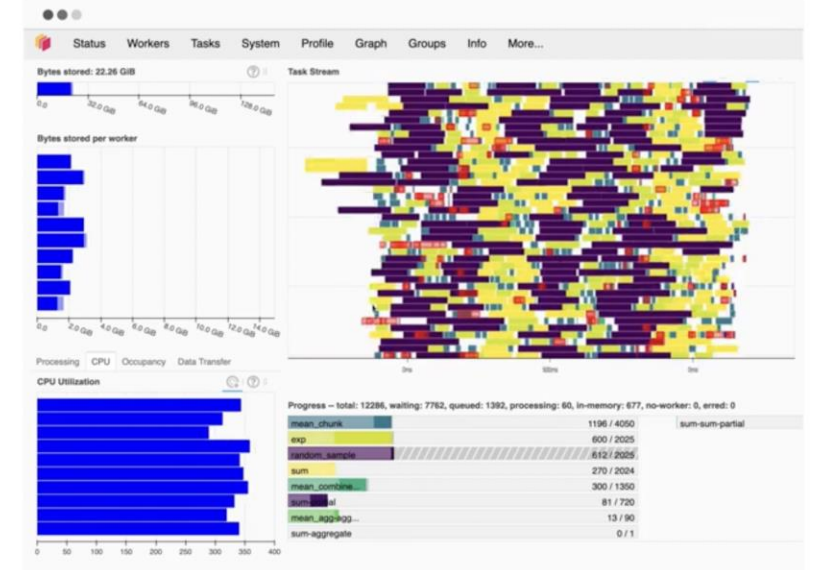


Python library for parallel and distributed computing

## Dask Arrays




## Dask Computation (Dask workers & Dask scheduler)



# What is Xarray and what is Dask?



N-D labeled arrays and datasets in Python

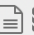



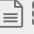



Supports chunking and  
lazy loading by  
wrapping Dask Arrays

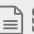
xarray.Dataset

► Dimensions: (lat: 25, time: 2920, lon: 53)


▼ Coordinates:

lat	(lat)	float32	75.0 72.5 70.0 ... 20.0 17.5 15.0		
lon	(lon)	float32	200.0 202.5 205.0 ... 327.5 33...		
time	(time)	datetime64[ns]	2013-01-01 ... 2014-12-31T18:0...		

▼ Data variables:

air	(time, lat, lon)	float32	dask.array<chunksize=(2920, ...		
-----	------------------	---------	---------------------------------	---	---

# Chunked data




latitude

(1000,300)

longitude




# Chunked data



latitude

(1000,300)

longitude



150

latitude

10 chunks  
(200,150)

longitude


200

# Dask and chunked data



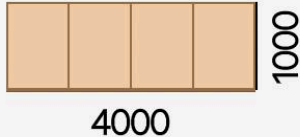
```
import dask.array as da
ones = da.ones(shape)
ones
```

	Array	Chunk
<b>Bytes</b>	30.52 MiB	30.52 MiB
<b>Shape</b>	(1000, 4000)	(1000, 4000)
<b>Count</b>	1 Tasks	1 Chunks
<b>Type</b>	float64	numpy.ndarray



```
chunk_shape = (1000, 1000)
ones = da.ones(shape, chunks=chunk_shape)
ones
```

	Array	Chunk
<b>Bytes</b>	30.52 MiB	7.63 MiB
<b>Shape</b>	(1000, 4000)	(1000, 1000)
<b>Count</b>	4 Tasks	4 Chunks
<b>Type</b>	float64	numpy.ndarray



## Activity 1:

How does chunking + parallel computing speed up computations?

## Activity 1:

How does chunking + parallel computing speed up computations?

**Lesson learned:** running computations on chunks in parallel speeds up the computation

Activity 2a:

How to best choose a chunk size?

Too ***small*** can be problematic!

Activity 2a:

How to best choose a chunk size?

Too ***small*** can be problematic!

**Lesson learned:** too small chunk sizes creates large overhead for scheduler and results in slow computation

Activity 2b:

How to best choose a chunk size?

Too ***big*** can also be problematic!

Activity 2b:

How to best choose a chunk size?

Too ***big*** can also be problematic!

**Lesson learned:** too big chunk sizes can cause memory errors and prevent the computation from finishing



## Activity 3:

How to decide which dimensions to chunk along?

## Activity 3:

How to decide which dimensions to chunk along?

**Lesson learned:** chunked dimensions must make sense for your computation

